SUDHAKAR CHUNDU

Staff DevOps / SRE & Cloud Infrastructure Leader

GPU / AI-ML Platforms · Platform Engineering · Multi-Cloud Architecture · Performance & HPC · FinOps

chundubabu@gmail.com | (513) 666-0099 | San Jose, CA | linkedin.com/in/babucs | sudhakarchundu.org | github.com/schundu007

PROFESSIONAL SUMMARY

Staff-level infrastructure leader with 18+ years of experience — 12+ years across cloud, platform, and HPC engineering built on 6 years of Linux and middleware foundations — spanning DevOps and Platform Engineering, Site Reliability Engineering, multi-cloud architecture, GPU / AI-ML infrastructure, and systems performance. Designs and operates large-scale platforms across bare-metal and multi-cloud (AWS, Azure, GCP, OCI): CI/CD and GitOps automation, Kubernetes (EKS/AKS/GKE and self-managed), Terraform IaC, 65-GPU training clusters, MLOps pipelines, and SRE practices (SLI/SLO, error budgets, incident management). Proven record of 99.97% uptime, $8M+ in annual cost savings, and 73% cloud cost reduction while leading a team of 6 engineers.

Currently Senior Data Platform Engineer at AT&T building secure Azure data platform infrastructure (ADLS Gen2 with Private Endpoints, Databricks, Delta Lake, Unity Catalog) for enterprise-scale AI/ML and analytics. Author of a vendor-neutral MLOps & LLMOps reference paper. Python is the primary language for automation and tooling. Certified AWS Solutions Architect Professional, Azure Solutions Architect Expert, Google Cloud Professional Cloud Architect, and CKA.

CORE TECHNICAL COMPETENCIES

Cloud & Multi-Cloud: AWS, Azure, GCP, OCI; hybrid and on-premise; Well-Architected landing zones, secure multi-account/subscription design, enterprise networking (VPC/VNet, ExpressRoute, Direct Connect, Private Endpoints, DNS), 15+ regions

IaC & GitOps: Terraform (5+ production modules, state management), ARM/Bicep, CloudFormation, Ansible, Pulumi, Helm, Kustomize; ArgoCD, FluxCD; policy-as-code (OPA, Checkov, tfsec)

Kubernetes & Containers: EKS, AKS, GKE, self-managed/on-prem; Docker, Helm, Istio service mesh, KEDA autoscaling, vcluster multi-tenancy, GPU scheduling, container security (Trivy, Falco)

CI/CD & Platform Engineering: Azure DevOps, GitHub Actions, Jenkins, GitLab CI; blue-green and canary; golden-path templates, self-service portals, internal CLIs, Backstage, infrastructure-as-product

GPU / AI-ML Infrastructure & MLOps: Multi-node GPU clusters (A100/V100/H100/T4, DGX), Slurm/Slinky, NVIDIA BCM/NGC/CUDA; PyTorch, Triton, TensorRT, vLLM, KServe, Ray Serve; MLflow, Kubeflow, Vertex AI, SageMaker, Azure ML; model quantization, edge AI (Jetson, K3s)

Performance & HPC: Linux profiling (perf, strace, ftrace, eBPF/bcc, Wireshark), flame graphs, GPU profiling (DCGM, Nsight, nvidia-smi), kernel tuning (cgroups, hugepages, NUMA), JVM/GC tuning, Databricks/Spark optimization

SRE & Observability: SLI/SLO/SLA, error budgets, incident management, blameless postmortems, chaos engineering, FMEA, toil reduction; Prometheus, Grafana, Datadog, ELK, Splunk, OpenTelemetry, custom exporters, AIOps

Data & Streaming: Databricks (Jobs, Compute, Unity Catalog), Delta Lake, Airflow; Apache Kafka, Flink, Spark, Hadoop/HDFS; exactly-once semantics, petabyte-scale pipelines, feature engineering

Identity, Security & Compliance: Entra ID, SailPoint, SSO, MFA, RBAC/ABAC, conditional access, workload identities, HashiCorp Vault; SOC2, HIPAA, HITRUST, FedRAMP, ISO 27001; Zero Trust, DevSecOps

FinOps: Cloud financial management, rightsizing, reserved-instance and spot strategies, storage tiering, showback/chargeback, vendor consolidation, ROI-driven investment

Programming: Python (primary — PyTorch, NumPy, APIs, automation), Bash, PowerShell, Java, Scala, SQL, Terraform HCL, Groovy, C/C++ (debugging)

PROFESSIONAL EXPERIENCE

Senior Data Platform Engineer — Azure Databricks & AI/ML Data Lake
Feb 2026 – Present

AT&T | San Jose, CA

- **Engineered** secure Azure ADLS Gen2 data lake with Private Endpoint connectivity, hub-and-spoke VNet integration, and Private DNS zones — eliminating public network exposure for petabyte-scale model training and enterprise data.
- **Operated** Databricks platform (Jobs, Compute, Unity Catalog) with Delta Lake ACID transactions for ML feature engineering, training pipelines, and multi-tenant data-science workspaces, including IBM watsonx Granite model integration.
- **Led** Wave 1 cloud-migration cutover and Magritte source onboarding under change management — integrating 50+ enterprise data sources with policy-as-code (OPA, Checkov, tfsec), automated schema discovery, and lineage.
- **Implemented** Entra ID identity federation, conditional access, and workload identities mapped to the Unity Catalog metastore for least-privilege model and feature access; automated service-principal lifecycle and secret rotation via Vault.

- **Tuned** Databricks/Delta Lake performance — Z-ordering, OPTIMIZE/VACUUM, file compaction, and Spark tuning — cutting ETL job runtimes 40% across business-critical pipelines.

Staff Platform Engineer / SRE & GPU Infrastructure Lead Oct 2023 – Jan 2026

Trackonomy Systems Inc | San Jose, CA

Led a 6-engineer infrastructure team (Cloud, DevOps, Security, Networking, Databases) serving 12+ enterprise clients across Pharma, Airlines, Government, Manufacturing, Healthcare, and IoT in 8 countries; owned reliability for a platform of 200+ microservices at 99.97% uptime.

GPU / AI-ML Infrastructure & MLOps

- **Built,** designed, and operated a Slurm-based 65-GPU training and compute platform across 8 bare-metal nodes (Slurm, Slinky, NVIDIA BCM) with preemption-safe scheduling, checkpointing, and fair-share policies — improving GPU utilization from 45% to 78% and cutting queue wait times 60%.
- **Architected** serverless GPU inference (Cloud Run, AWS Lambda containers) and multi-platform pipelines (NVIDIA CUDA/TensorRT, Apple MLX, CPU fallback) across 15+ edge locations — reducing inference costs 65% while serving 5M+ daily predictions.
- **Engineered** automated ML pipelines with ArgoCD/GitOps, experiment tracking, and blue-green deployments for 30+ production models — cutting deployment time from days to 15 minutes; secured the GenAI/LLM platform against prompt injection and data exfiltration.

SRE, Reliability & Observability

- **Ran** L3 SRE operations — RCA, incident management, blameless postmortems, SLI/SLO and error-budget tracking, and 24x7 on-call (<15-min response) — reducing MTTR 60%; applied AIOps for predictive monitoring and anomaly detection, cutting incidents 65%.
- **Unified** 5 monitoring platforms into Datadog and built a Prometheus/Grafana/OpenTelemetry stack with custom DCGM/GPU exporters and distributed tracing; executed chaos engineering, FMEA, and DR validation achieving RPO <1h / RTO <4h.

Platform, CI/CD & FinOps

- **Built** CI/CD for 100+ applications with GitOps (ArgoCD, FluxCD), GitHub Actions, Jenkins, and Azure DevOps — enabling 50+ daily zero-downtime deployments and reducing production issues 90%; delivered golden-path templates and self-service tooling for 50+ engineers.
- **Redesigned** multi-cloud architecture (AWS, Azure, GCP, OCI), reducing Azure resources 87% (8.7K→1.1K), and drove 73% cloud cost reduction ($10M→$2.7M/yr) — delivering $8M+ in annual savings through GPU optimization, fair-share scheduling, and reserved-instance planning.

Senior SRE / Cloud Architect — OSDU Data Platform Mar 2022 – Oct 2023

AWS OSDU via Tecra Systems (Wipro, contract) | Hyderabad, India

- **Architected** the OSDU data platform on GPU-accelerated Kubernetes (EKS/Fargate) processing exabytes of seismic data with Spark, Kafka, and Hadoop/HDFS for ExxonMobil, Chevron, BP, and Shell — ensuring 99.95%–99.99% availability for ML inference.
- **Created** 5+ reusable Terraform modules (VPC, EKS, RDS, DynamoDB, S3, IAM) on Well-Architected principles with multi-tenant AWS Landing Zones (Organizations, SCPs); implemented GitOps (ArgoCD/FluxCD) and KEDA — reducing deployment time 80%.
- **Built** Prometheus/Grafana/ELK observability and designed DR (RPO <1h, RTO <4h); led 24x7 on-call and MLOps (model versioning, automated retraining, blue-green for 30+ models via MLflow).

Cloud Infrastructure Engineer — Azure Platform Engineering Feb 2020 – Mar 2022

Microsoft Azure OSDU Platform (Wipro, full-time) | Hyderabad, India

- **Engineered** scalable infrastructure on Azure AKS with GPU node pools and ExpressRoute hybrid connectivity; deployed 15+ microservices with Istio service mesh (mTLS, distributed tracing) achieving 99.9% availability.
- **Built** Azure DevOps pipelines with Terraform and Bicep IaC, approval gates, and security scanning (Checkov, tfsec) — improving deployment speed 80%; delivered compliance-as-code achieving SOC 2, HIPAA, and ISO 27001.

Cloud Architect — Healthcare Infrastructure Aug 2018 – Feb 2020

NTT Data Services | Hyderabad, India

- **Led** cloud modernization for HIPAA/HITRUST-regulated applications to AWS; implemented GPU-enabled EKS clusters and compliant ML pipelines with end-to-end encryption, audit logging, and secure model serving; managed a $6M infrastructure budget.
- **Built** Jenkins/Ansible CI/CD and DevSecOps (HashiCorp Vault, SonarQube, OWASP); designed DR with business-continuity planning.

Multi-Cloud Architect / Senior Infrastructure Engineer May 2007 – Jun 2018

Tata Consultancy Services — CNA, PwC, Verizon, Owens Corning, ASML | India & USA

- **Progressive** 11-year career across Fortune 500 clients: Linux/Unix administration → cloud architecture → infrastructure leadership, including 12+ facility buildouts (branch offices, call centers, data centers).
- **Served as** Senior Linux Administrator, ASML (2009–2011) — administered large-scale RHEL/Linux server fleets for mission-critical semiconductor engineering environments; automated provisioning, patching, and monitoring; tuned kernel and storage for high-availability compute workloads with 24x7 on-call support.
- **Delivered** Middleware DevOps for PwC — WebSphere and WebLogic administration, automation, and deployment pipelines on Azure and AWS; managed application-server clustering, JVM tuning, and zero-downtime middleware releases.
- **Led** cloud migrations to AWS, Azure, and OpenStack; pioneered Docker/Kubernetes adoption (2013–14); implemented CI/CD with Jenkins, Ansible, and Terraform; deep Linux administration and middleware operations across enterprise workloads with 24x7 on-call.

KEY ACHIEVEMENTS & METRICS

Reliability & Scale: 99.97% uptime · 65-GPU multi-node cluster · 200+ microservices · exabyte-scale data · 15+ regions · 5M+ daily predictions

Cost & Efficiency: $8M+ annual savings · 73% cloud cost reduction · 87% infrastructure-sprawl elimination · 65% inference cost reduction

Velocity: 95% deployment-time reduction (days → 15 min) · 50+ daily zero-downtime deployments · 80% faster provisioning

Operational Excellence: 60% MTTR improvement · 65% incident reduction via AIOps · <15-min on-call response · 85% manual-ops reduction

Performance: GPU utilization 45% → 78% · queue wait times down 60% · ETL runtimes down 40% · build times down 60%

PORTFOLIO, OPEN SOURCE & PUBLICATIONS

Portfolio: sudhakarchundu.org · **Blog:** sudhakarchundu.org/blog · **Medium:** schundu.medium.com

Open Source: OSDU Forum — infra-azure-provisioning (Terraform + AKS automation, 2,317+ commits) and terraform-deployment-aws; GitHub: github.com/schundu007

Publication: Author, vendor-neutral MLOps & LLMOps reference paper — orcid.org/0009-0001-5038-5936

CERTIFICATIONS

Cloud: AWS Solutions Architect Professional | Azure Solutions Architect Expert | Google Cloud Professional Cloud Architect

DevOps & Data: Certified Kubernetes Administrator (CKA) | HashiCorp Terraform Associate | Databricks Certified Data Engineer

Enterprise: TOGAF 9 Foundation | ITIL v4 Foundation | VMware VCP | Cisco CCNP

EDUCATION

Bachelor of Engineering, Mechanical Engineering — Acharya Nagarjuna University, 2005